

# Machine Learning 1

Katherine Wong (A16162648)

10/21/2021


First up is clustering methods

## Kmeans clustering

The function in base R to do Kmeans clustering is called 'kmeans'.

First make up some data where we now what the answer should be:

```
tmp <- c(rnorm(30, -3), rnorm(30, 3))  
x <- cbind(x = tmp, y = rev(tmp))  
plot(x)
```



Q. Can we use kmeans() to cluster this data setting k 2 and nstart to 20?

```
km <- kmeans(x, centers = 2, nstart = 20)
km

## K-means clustering with 2 clusters of sizes 30, 30
##
## Cluster means:
##           x           y
## 1  3.004976 -3.160746
## 2 -3.160746  3.004976
##
## Clustering vector:
## [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
## [39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
##
## Within cluster sum of squares by cluster:
## [1] 77.17282 77.17282
## (between_SS / total_SS =  88.1 %)
##
## Available components:
##
## [1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
## [6] "betweenss"    "size"        "iter"        "ifault"
```

Q. How many points are in each cluster?

```
km$size

## [1] 30 30
```

Q. What 'component' of your result object details cluster assignment / membership?

```
km$cluster

## [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
## [39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


Q. What 'component' of your result object details cluster center?

```
km$centers

##           x           y
## 1  3.004976 -3.160746
## 2 -3.160746  3.004976
```

Q. Plot x colored by the kmeans cluster assignment and add cluster centers as blue points

```
plot(x, col = km$cluster)
points(km$centers, col="blue", pch = 15, cex = 2)
```



## Hierarchical Clustering

A big limitation with k-means is that we have to tell it K (the number of clusters we want).

Analyze this same data with `hclust()`.


Demonstrate the use of `dist()`, `hclust()`, `plot()` and `cutree()` functions to do clustering, Generate dendrograms and return cluster assignment membership vector...

```
hc <- hclust(dist(x))
hc

##
## Call:
## hclust(d = dist(x))
##
## Cluster method   : complete
## Distance         : euclidean
## Number of objects: 60
```

There is a plot method for `hclust` result objects. Let's see it.

```
plot(hc)
```



To get our cluster membership vector we have to do a wee bit more work. We have to “cut” the tree where we think it makes sense. For this we use the ‘cutree()’ function.

```
cutree(hc, h=6)
```


```
## [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
## [39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

You can also call ‘cutree()’ setting  $k$  = the number of groups / clusters you want.

```
grps <- cutree(hc, k = 2)
```

Make our results plot

```
plot(x, col = grps)
```



## Princial Component Analysis (PCA)

Read data on food stuffs from

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
## [1] 17  5
```

There are 17 rows and 5 columns. You can use the `dim()` function to get the rows and columns or call `nrow()` to get rows and `ncol()` to get columns separately.

## Checking data

Preview the first 6 rows

```
head(x)
```

```
##           X England Wales Scotland N.Ireland
## 1      Cheese      105   103      103       66
## 2 Carcass_meat     245   227      242      267
## 3   Other_meat     685   803      750      586
## 4         Fish     147   160      122       93
## 5 Fats_and_oils     193   235      184      209
## 6       Sugars     156   175      147      139
```

There is an extra col as reported by the `dim()` function.

```
# Note how the minus indexing works
#rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

```
##   England Wales Scotland N.Ireland
## 1     105   103      103       66
## 2     245   227      242      267
## 3     685   803      750      586
## 4     147   160      122       93
## 5     193   235      184      209
## 6     156   175      147      139
```

```
# Or do this when loading data
# x <- read.csv(url, row.names=1)
# head(x)
```

Check dimensions again

```
dim(x)
```

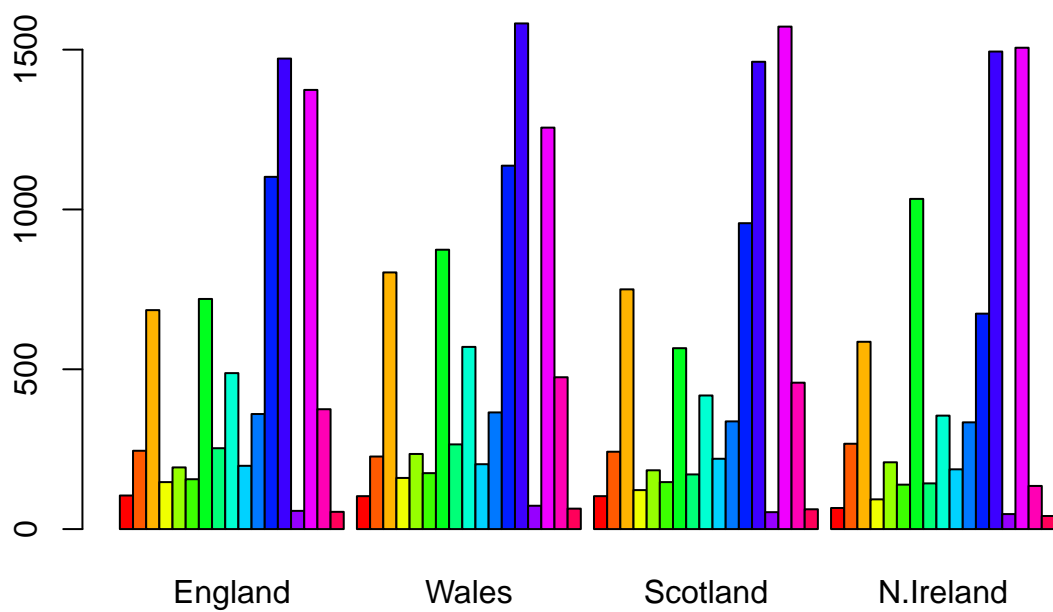
```
## [1] 17  4
```

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

Loading it and setting `row.names` to the first column is the preferred method. For the first method, if you call `x <- x[,-1]` multiple times, the columns start disappearing.


## Spotting major differences and trends

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



A: If you change `beside=T` to `beside=F`, you can achieve this barplot.

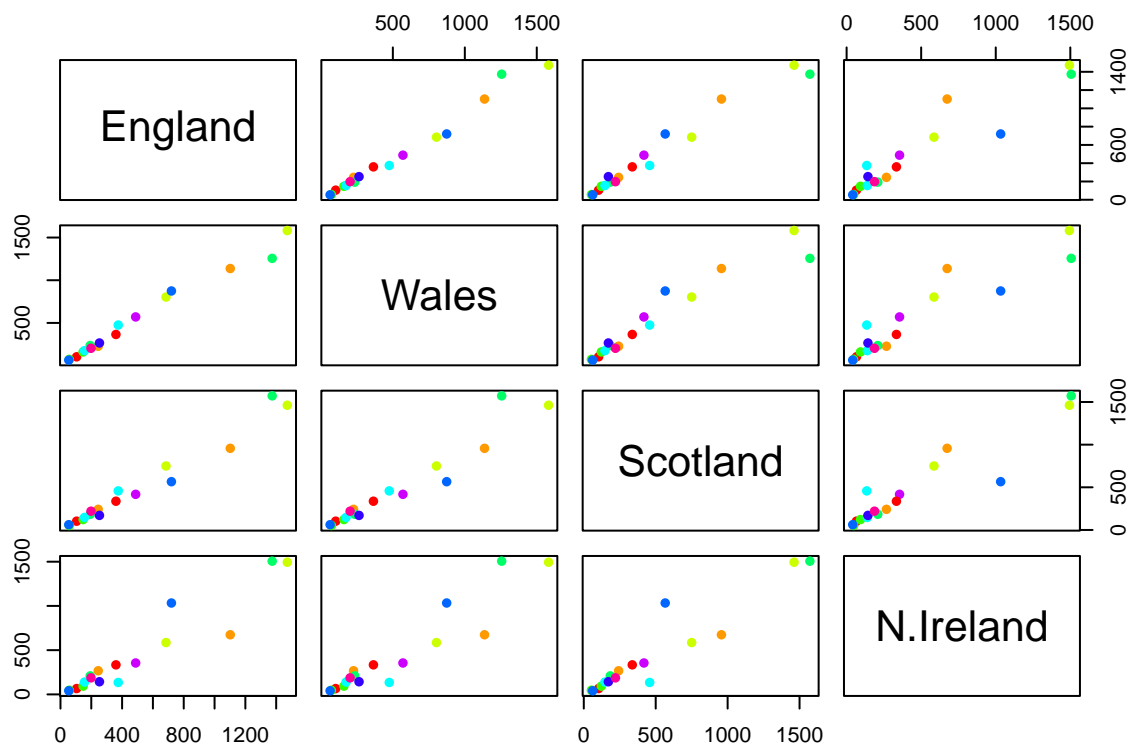
The stacked barplot is not helpful.

Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

Each row is the y axis (for example, first row is england) and column is the x axis (for example, first column is england). If it lies on the diagonal then the two countries have similar food consumption for all food types. If it deviates from the diagonal, then the countries don't have similar food consumption numbers.

```
pairs(x, col=rainbow(10), pch=16)
```





Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

N.Ireland has blue dots higher on the y axis than any other country. The orange dot also seems to be lower on the y axis compared to the other UK countries.

## PCA to the rescue!

The main function in base R for PCA is 'prcomp()' This wants the transpose of our data

```
pca <- prcomp(t(x))
summary(pca)
```

```
## Importance of components:
```


```
##          PC1      PC2      PC3      PC4
## Standard deviation 324.1502 212.7478 73.87622 4.189e-14
## Proportion of Variance 0.6744 0.2905 0.03503 0.000e+00
## Cumulative Proportion 0.6744 0.9650 1.00000 1.000e+00
```

```
attributes(pca)
```

```
## $names
```


```
## [1] "sdev"      "rotation" "center"   "scale"    "x"
##
## $class
## [1] "prcomp"
```

```
plot( pca$x[,1], pca$x[,2])
```




> Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```



> Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at the start of this document

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x), col = c("orange", "red", "blue", "dark green"))
```



```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```


```
## [1] 67 29 4 0
```

or the second row here...

```
## or the second row here...
z <- summary(pca)
z$importance
```


```
##              PC1      PC2      PC3      PC4
## Standard deviation 324.15019 212.74780 73.87622 4.188568e-14
## Proportion of Variance 0.67444 0.29052 0.03503 0.000000e+00
## Cumulative Proportion 0.67444 0.96497 1.00000 1.000000e+00
```

```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```



## Digging Deeper (variable loadings)


```
## Lets focus on PC1 as it accounts for > 90% of variance
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,1], las=2 )
```



Q9: Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

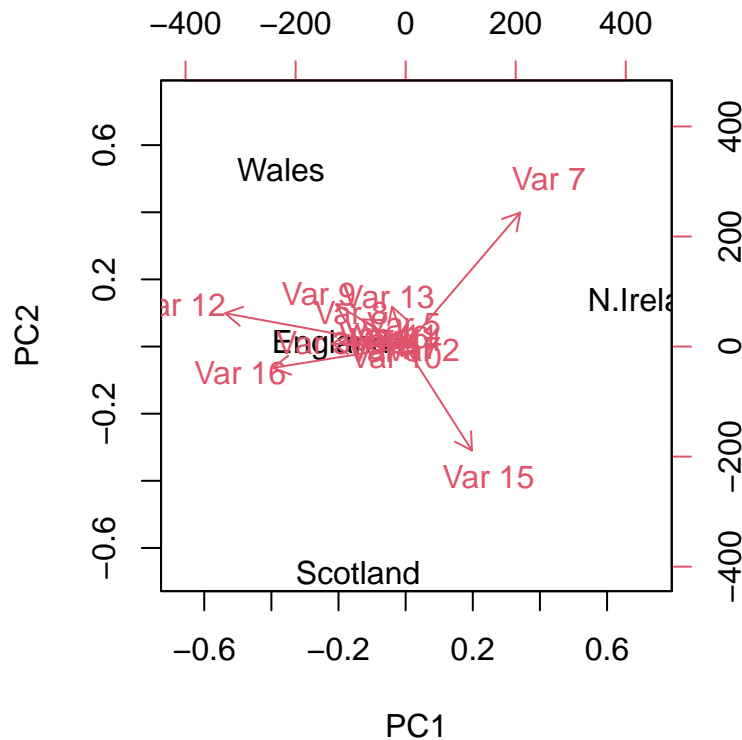
Fresh potatoes and soft drinks are positive. Negative is other meat, fresh fruit and alcoholic drinks

```
## Lets focus on PC2
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```



The inbuilt `biplot()` can be useful for small datasets

```
biplot(pca)
```



#PCA of RNA Seq Data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

```
##          wt1 wt2  wt3  wt4 wt5 ko1 ko2 ko3 ko4 ko5
## gene1    439 458  408  429 420  90  88  86  90  93
## gene2    219 200  204  210 187 427 423 434 433 426
## gene3   1006 989 1030 1017 973 252 237 238 226 210
## gene4    783 792  829  856 760 849 856 835 885 894
## gene5    181 249  204  244 225 277 305 272 270 279
## gene6    460 502  491  491 493 612 594 577 618 638
```

Q10: How many genes and samples are in this data set?

```
dim(rna.data)
```


```
## [1] 100  10
```

100 genes, and 10 samples.

```
## Again we have to take the transpose of our data
pca <- prcomp(t(rna.data), scale=TRUE)
```



```
## Simple un polished plot of pc1 and pc2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2")
```



```
summary(pca)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation    9.6237  1.5198  1.05787  1.05203  0.88062  0.82545  0.80111
## Proportion of Variance 0.9262  0.0231  0.01119  0.01107  0.00775  0.00681  0.00642
## Cumulative Proportion 0.9262  0.9493  0.96045  0.97152  0.97928  0.98609  0.99251
##              PC8      PC9      PC10
## Standard deviation    0.62065  0.60342  3.348e-15
## Proportion of Variance 0.00385  0.00364  0.000e+00
## Cumulative Proportion 0.99636  1.00000  1.000e+00
```

```
plot(pca, main="Quick scree plot")
```

## Quick scree plot



```
## Variance captured per PC
```

```
pca.var <- pca$sdev^2
```

```
## Percent variance is often more informative to look at
```


```
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
```

```
pca.var.per
```

```
## [1] 92.6  2.3  1.1  1.1  0.8  0.7  0.6  0.4  0.4  0.0
```

```
barplot(pca.var.per, main="Scree Plot",  
        names.arg = paste0("PC", 1:10),  
        xlab="Principal Component", ylab="Percent Variation")
```


## Scree Plot



```
## A vector of colors for wt and ko samples
colvec <- colnames(rna.data)
colvec[grep("wt", colvec)] <- "red"
colvec[grep("ko", colvec)] <- "blue"

plot(pca$x[,1], pca$x[,2], col=colvec, pch=16,
     xlab=paste0("PC1 (", pca.var.per[1], "%)"),
     ylab=paste0("PC2 (", pca.var.per[2], "%)"))

text(pca$x[,1], pca$x[,2], labels = colnames(rna.data), pos=c(rep(4,5), rep(2,5)))
```

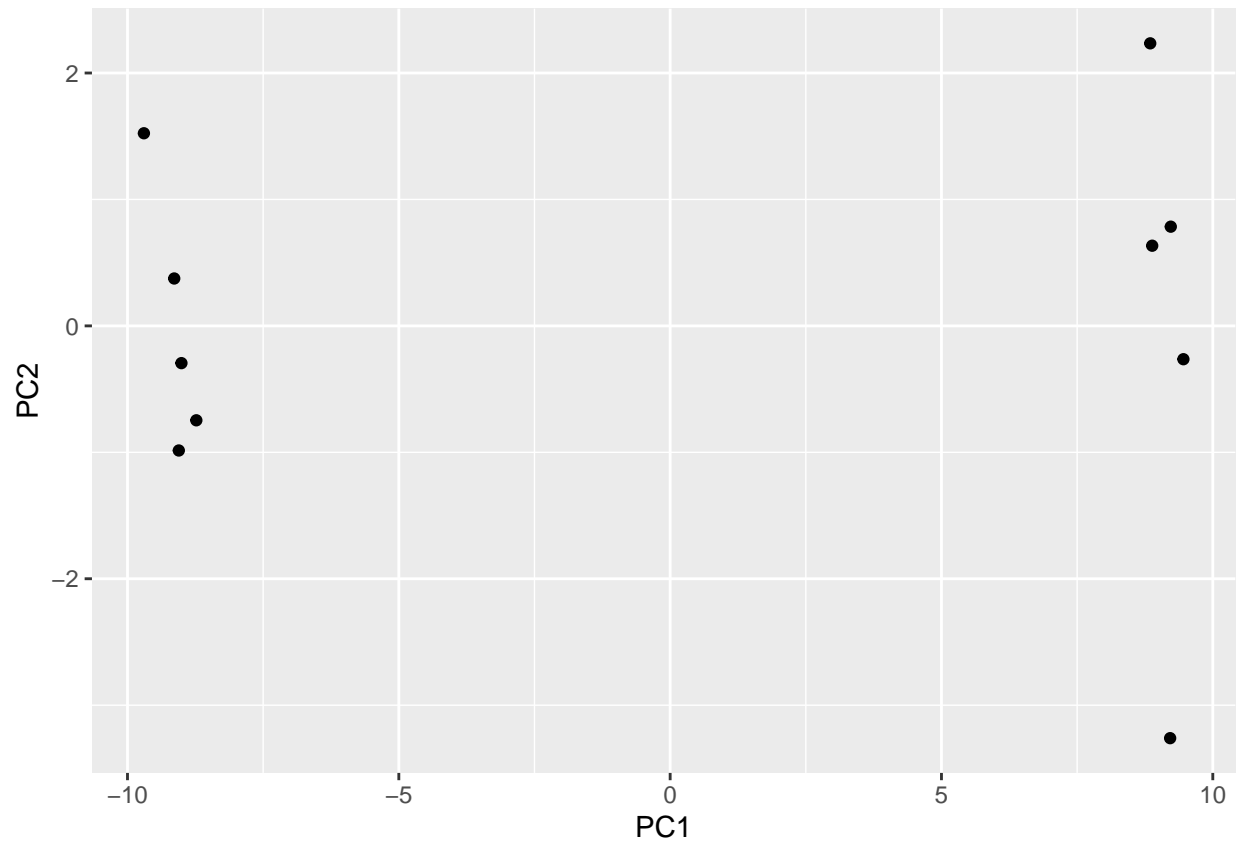


# Using ggplot

```
library(ggplot2)

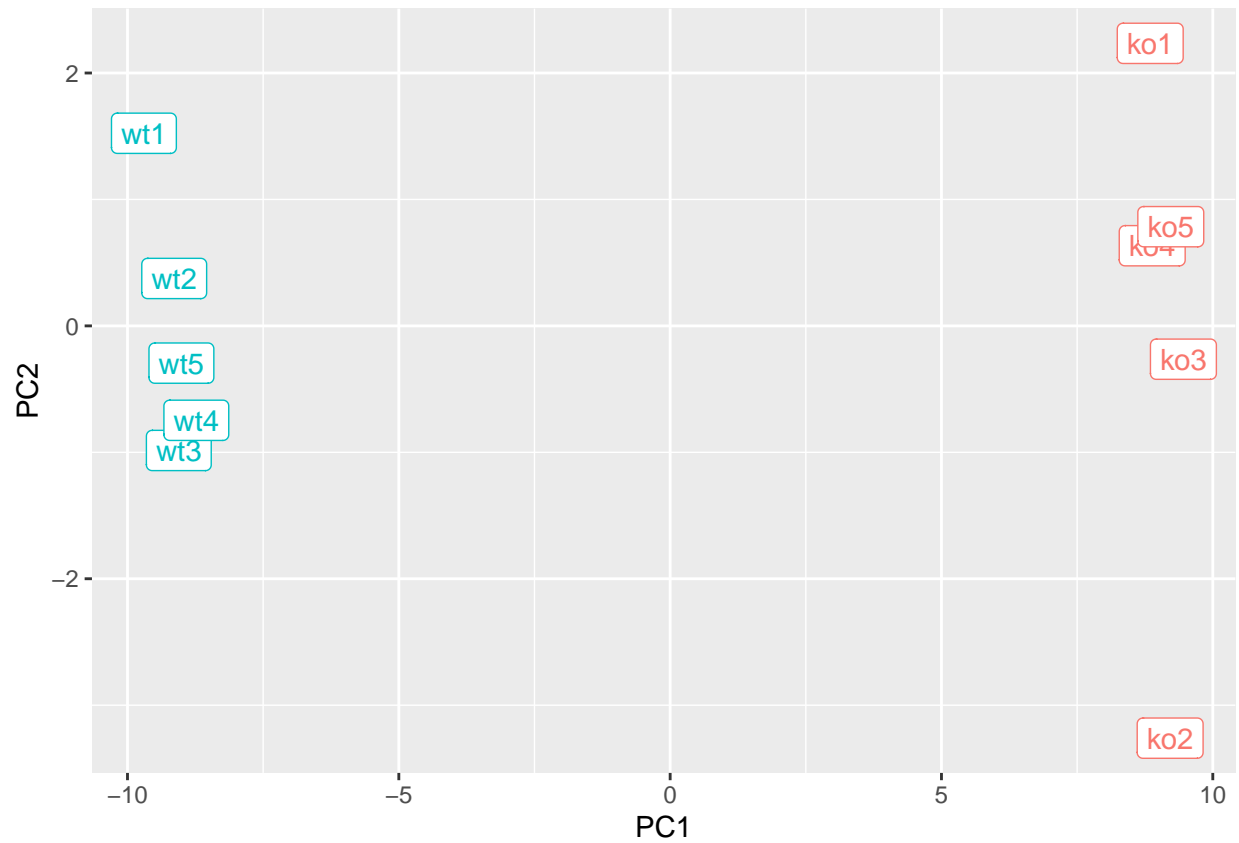
df <- as.data.frame(pca$x)

# Our first basic plot
ggplot(df) +
  aes(PC1, PC2) +
  geom_point()
```



```
# Add a 'wt' and 'ko' "condition" column
df$samples <- colnames(rna.data)
df$condition <- substr(colnames(rna.data),1,2)


p <- ggplot(df) +
  aes(PC1, PC2, label=samples, col=condition) +
  geom_label(show.legend = FALSE)
p
```



```
p + labs(title="PCA of RNASeq Data",
  subtitle = "PC1 clealy seperates wild-type from knock-out samples",
  x=paste0("PC1 (", pca.var.per[1], "%)"),
  y=paste0("PC2 (", pca.var.per[2], "%)"),
  caption="BIMM143 example data") +
  theme_bw()
```

## PCA of RNASeq Data

PC1 clearly separates wild-type from knock-out samples



BIMM143 example data